

Week 13: MST (Prim) and SSSP (Dijkstra)

Agenda:

- ▶ Minimum Spanning Trees (Cont'd): Prim's Algorithm
- ▶ Single-Source Shortest Paths: Dijkstra's Algorithm

Reading:

- ▶ CLRS: 624-642 (covers Prim's)
- ▶ CLRS: 643-664 (covers Dijkstra's)

Recall: Minimum Spanning Tree (MST) problem

- ▶ Input: edge-weighted undirected graph
- ▶ Notions:
 - ▶ subgraph: $G' = (V, E')$, where $E' \subset E$, forest = acyclic graph, trees
 - ▶ spanning subgraph: subgraph including all the vertices
 - ▶ spanning tree: spanning subgraph which is a tree — acyclic connected must have exactly $n - 1$ edges
e.g., BFS/DFS-tree is a spanning tree of the graph
 - ▶ minimum spanning tree: sum of weights on tree edges is minimal
- ▶ **The MST Problem: Find a minimum spanning tree (MST) for the input graph.**
 - ▶ ... there could be more than one ..., e.g., all weights are the same, both BFS/DFS produce a MST ...
- ▶ Important for:
 - ▶ Min-cost set of edges that we need so that all vertices can reach one another
 - ▶ Learning value: a canonical example for greedy algorithms.
 - ▶ Useful info derived from the MST algorithms...
- ▶ The Minimum Spanning Forest problem: If the given graph is not necessarily connected: find MST for each CC.

Recall: Greedy algorithms and MST problem

- ▶ Greedy algorithms:
 - ▶ greedy — each step makes the best choice (locally minimum)
 - ▶ and don't look back
 - ▶ Optimal substructure: an optimal solution to the original problem contains within it optimal solutions to subproblems.
 T is MST for $G = (V, E) \Rightarrow$ for any $U \subset V$ where $T[U]$ is connected, $T[U]$ is a min-spanning tree.
- ▶ The general MST algorithm outline:
 1. A is a set of “safe” edges: they are contained in some MST T
 2. $A = \emptyset$ initially
 3. **while** ($|A| < n - 1$) **do**: find a safe edge $e = (u, v)$ and set $A = A \cup \{e\}$.

What happens when we halt?
- ▶ Two greedy solutions
 - ▶ Prim's Algorithm (Actually: Prim + Dijkstra + Boruvka)
 Grow T vertex-wise: A is always a MST on some $S \subset V$
 - ▶ Kruskal's (Actually: Kruskal + Boruvka)
 Grow T edge-wise: A is always a minimal set of edges without a cycle (forest)

Prim's algorithm for the MST problem:


- ▶ Input: an edge-weighted (simple, undirected, connected) graph (positive weights)
- ▶ Output: an MST
- ▶ Idea:
 - ▶ Suppose we have already an MST A spanning subset S of vertices (Initially: $S = \text{a single vertex}$, $A = \emptyset$)
 - ▶ Grow A to span one more vertex $v \in \bar{S} = V \setminus S$ by adding a single edge (u, v) for some $u \in S$ and $v \notin S$.
 - ▶ Which edge to pick? Greedy! min-weight edge from all possible edges crossing the (S, \bar{S}) cut.
- ▶ First sketch:

```


procedure primMST( $G$ )      ** $G = (V, E)$ 
 $S = \{s\}$                   **for any start vertex  $s$ 
 $A = \emptyset$ 
while ( $|S| < |V|$ ) do
    find a minimum weight edge  $e = (u, v)$ :  $u \in S$  and  $v \notin S$ 
     $S = S \cup \{v\}$ 
     $A = A \cup \{e\}$ 
return  $A$ 
  
```

Prim's algorithm for the MST problem — an example:

- ▶ Input graph G :



- ▶ $\text{primMST}(G, w, 1)$ returns:




- ▶ First we prove correctness of Prim's algorithm
- ▶ Then we improve the naïve algorithm to reduce runtime
Not surprisingly, finding the min-edge quickly is going to be useful (heaps!)

Prim's algorithm for the MST problem — correctness:

- ▶ In the proof we show the substitution property.
- ▶ At any point, suppose $A \subset T$ for some MST T , and we show that there exists a MST T' that contains $A \cup \{e\}$.
- ▶ If $e \in T$, we are done.
- ▶ Since $u \in S$ whereas $v \notin S$, so on the unique $u \rightarrow v$ path there has to be some edge $e' = (u', v')$ in T that crosses the (S, \bar{S}) -cut. Clearly, $w(e) \leq w(e')$.
- ▶ We argue $T' = T \setminus \{e'\} \cup \{e\}$ is spanning V ; and as it has $n - 1$ edge it has to a spanning tree, with cost $\leq w(T)$.
- ▶ In fact, it is enough to argue u' and v' remain connected:
Any $x \rightarrow y$ path on T either goes through e' or not. In the latter case, the $x \rightarrow y$ path remains in the T' ; in the former case — use the new $u' \rightarrow v'$ path to connect x with y .
- ▶ So why do u' and v' remain connected? Well, in T there's a $u \rightarrow v$ path that e' was a part of. So u is connected to u' and v' is connected to v . So, $u' \rightarrow \underbrace{u, v}_e \rightarrow v'$ is a path connected u' and v' in T' . ■

Prim's algorithm for the MST problem — faster implementation:



► Example: input graph G :



► $\text{primMST}(G, w, 1)$ returns:

► $\text{primMST}(G, w, 1)$: an intermediate tree

Already picked black edges spanning 1,5,3; what are the **candidate edges**?



Prim's algorithm for the MST problem — faster implementation:

Idea:

For each node $\notin S$ — keep track of the min edge that connects it to S .

Update the information only for the neighbors of the node that is currently being added to S .

Uses a priority queue Q on \bar{S} (so S is implicit — all the nodes **not** in Q)

Pseudocode:

procedure primMST(G)

 for each $v \in V(G)$ do

$v.key = \infty$

$v.predec = \text{NIL}$

$s.key = 0$

 Initialize a min-priority-queue Q on V using key

 while ($Q \neq \emptyset$) do

$u = \text{ExtractMin}(Q)$

 ** u dequeued first


 foreach v neighbor of u do

 if ($v \in Q$ and $w(u, v) < v.key$) then

$v.predec = u$

 decrease-key($Q, v, w(u, v)$) ** $v.key$ is now $w(u, v)$

~~✗~~




maintain $\{Q, \bar{Q}\}$ on V, E
where $V, \bar{Q} = \min\{w(u, v) \mid u \in S\}$
invariant

Prim's algorithm for the MST problem — faster implementation:


- ▶ Analysis of the improved algorithm:
 - ▶ correctness (almost done — need to prove that $\text{ExtractMin}(Q)$ does extract the minimum weight edge that crosses the (Q, \bar{Q}) -cut)
 - ▶ running time: $\Theta\left(n + \sum_{u \in V} \left((1 + \deg(u)) \cdot \log n\right)\right)$
 - so: $\Theta((n + m) \log n) = \Theta(m \log(n))$ — adjacency list graph representation for a connected graph $m \geq n - 1$.
 - or $\Theta(n^2 + m \log(n))$ in the adjacency matrix representation.
- ▶ Using more sophisticated data-structure (Fibonacci heap), runtime of Prim is reduced to $O(n \log(n) + m)$.

What does Prim Teach Us?

- ▶ Theorem: Let T be a spanning tree of G .
Then T is a MST iff each edge is the min-edge crossing the cut it induces.
- ▶ I.e., take T , and some $e \in T$. Removing e from T disconnects T and creates two components $C_1, C_2 = V \setminus C_1$.
Then the claim is that e is a minimal edge out of all the edges the cross the (C_1, C_2) -cut.
Moreover, this is true for all edges $e \in T$.



MST:



- ▶ E.g.:
min-weight of edge crossing $(\{1, 5, 6\}, \{2, 3, 4, 7\})$ -cut = 8
min-weight of edge crossing $(\{1, 2, 3, 5, 6\}, \{4, 7\})$ -cut = 13
- ▶ Proof. \Rightarrow If T is a MST, pick any e . For contradiction, assume e isn't a minimum edge crossing the cut it induces – replace e with an edge of strictly smaller weight. The resulting graph is a spanning tree (has $n - 1$ edges and spans V) with cost strictly smaller than T . Contradiction.
- ▶ \Leftarrow If all edges in T satisfy this property — how do we prove it is a MST?
- ▶ Use Prim!

What does Prim Teach Us?


- ▶ Theorem: Let T be a spanning tree of G .
Then T is a MST iff each edge is the min-edge crossing the cut it induces.
- ▶ = If all edges in T satisfy this property — how do we prove it is a MST?
- ▶ Use Prim!
Claim: there's an instantiation of Prim that builds T , starting from s .
In other words: we can run Prim and maintain the invariant that $T[S]$ is a single connected component ($T[S]$ is a spanning tree of S).
- ▶ Proof by induction on S . Clearly true for $|S| = 1$.
- ▶ The induction step:
 - ▶ Suppose in the transition from S to $S \cup \{v\}$ Prim picks an edge $e = (u, v)$ of weight w . Let $e_1 = (a_1, b_1), e_2 = (a_2, b_2), \dots, e_k = (a_k, b_k)$ be all the edges in T with one vertex (a_i) in S and one vertex (b_i) in \bar{S} .
 - ▶ Since Prim picks the min-edge connecting S with \bar{S} , we have $w(e) \leq w(e_1), w(e) \leq w(e_2), \dots, w(e) \leq w(e_k)$.
 - ▶ ASOC $w(e) < \min\{w(e_1), w(e_2), \dots, w(e_k)\}$ — the weight of the edge chose by Prim is strictly smaller than all of the edges in T that leave S .
 - ▶ Look at the path $v \rightarrow u$ on T . It starts at \bar{S} and edges at S , so it must use some edge e_j . So the removal of e_j separates u from v .
 - ▶ Hence e_j isn't the min-edge that separates the cut (e also crosses the same cut). Contradiction!
 - ▶ Thus $w(e) = \min\{w(e_1), \dots, w(e_k)\} = w(e_j)$.
 - ▶ Instantiate the priority-queue of Prim to pick b_j rather than v (both have the same *key*, so break ties in favor of b_j rather than v).

Shortest path problems:

- ▶ BFS recall: outputs every s -to- v shortest path
 - ▶ s — start vertex
 - ▶ v — reachable vertex from s (residing in a same connected component)
 - ▶ shortest — # edges
 - ▶ running time $\Theta(n + m)$

- ▶ But what if the edges of the graph have weights?

In this case, shortest-path in terms of #edges isn't shortest weighted path.




E.g. shortest weighted-path distance between 1 and 2 is 8.

- ▶ The weight of a path = sum of weights on edges on path.
Note: if there is no path between two nodes, the distance is set to ∞ ...
- ▶ The SHORTEST PATH problem: find the shortest path in an edge-weighted graph connecting s and t .
- ▶ Turns out, shortest-path has a few properties that allow us to infer in the process **all** shortest-paths from a node s to any other node in the graph.
So we study the SINGLE-SOURCE SHORTEST PATH (SSSP) problem:
Given an edge-weighted graph G and a source s , find out for each vertex $v \in V(G)$ a shortest paths from s to v .
- ▶ Variants:
 - ▶ edge weights: non-negative vs. arbitrary weights

Shortest Paths

- ▶ A shortest path from u to v : out of all $u \rightarrow v$ paths, it is a path of minimal weight. (Can be more than one.)
- ▶ But a shortest path must satisfy **subpath optimality**:
if $(u_0, u_1, \dots, \underbrace{u_i, \dots, u_j}_{\text{optimal}}, \dots, u_k)$ is a shortest $u_0 \rightarrow u_k$ path, then $(u_i, u_{i+1}, \dots, u_j)$ is a shortest $u_i \rightarrow u_j$ path.
- ▶ we denote $d(u, v)$ as the shortest path length from u to v
- ▶ Shortest paths remain well-defined if there are negative weight directed



edges... (why not undirected?)

- ▶ ...as long as there isn't a negative weight path from a node to itself.
- ▶ So we assume **no negative cycles**.
- ▶ Shortest path distances — d is a metric:
 - ▶ For any u , we have $d(u, u) = 0$
 - ▶ For any u, v , we have $d(u, v) = d(v, u)$ (if G is undirected)
 - ▶ For any u, v, w , we have $d(u, w) \leq d(u, v) + d(v, w)$.

Common Outline of Single Source Shortest Path Algorithms

- ▶ Our shortest paths algorithm starts at a source s .
- ▶ It maintains a *dist* attribute for each vertex that will serve as the estimation of the shortest-path distance.
- ▶ During the execution of the algorithm, we **always** have $u.dist \geq d(s, u)$.
- ▶ We start with `init()`: set $s.dist \leftarrow 0$ and $u.dist \leftarrow \infty$ for any $u \neq s$.
- ▶ We only update the *dist* attribute by Relaxing

update \approx

procedure relax(u, v) **** u, v two adjacent nodes**


if ($v.dist > u.dist + w(u, v)$) then

$v.dist \leftarrow u.dist + w(u, v)$ \rightarrow update shortest path

- ▶ Claim: any algorithm that starts with `init()` and only updates *dist* using `relax()` must always satisfy $v.dist \geq d(s, v)$ for any v .
- ▶ Proof: by induction on the number of times `relax()` is invoked.
Base case: invoked 0 times — claim holds through `init()`.
Induction step: If `relax(u, v)` doesn't change $v.dist$ we are done.
Otherwise, we now have

$$v.dist = u.dist + w(u, v) \geq d(s, u) + w(u, v) \geq d(s, v)$$

- ▶ Corollary: Since `relax()` cannot increase $v.dist$, so if and when we set $v.dist = d(s, v)$ we keep $v.dist$ unchanged.



by induction, $v.d \geq d(s, v)$

by triangle inequality
 $d(s, v) \leq d(s, u) + w(u, v)$
 $\Rightarrow v.d + w(u, v)$

Dijkstra's SSSP algorithm:

- ▶ For graphs with non-negative weights (both directed and undirected)
- ▶ Idea in Dijkstra's algorithm:
 - ▶ Maintains a set S of vertices for which we know the shortest path. Initially, $S = \{s\}$, at the end: $S = V$.
 - ▶ Which vertex from $\bar{S} = V \setminus S$ should we pick?
 - ▶ Dijkstra: the greedy solution — the node in \bar{S} with minimum $dist$.


```

▶ procedure dijkstra( $G, w, s$ )           ** $G = (V, E)$ ,  $w$  =weights
  foreach  $v$  do                          **initialization
     $v.dist \leftarrow \infty$ 
     $v.predec \leftarrow \text{NIL}$ 
   $s.dist \leftarrow 0$ 
  Build Min-Priority-Queue  $Q$  on all nodes, key =  $dist$ 
    ** nodes in  $Q$  are nodes we are not yet sure about
    ** namely  $Q$  holds nodes in  $\bar{S}$ 
  while ( $Q \neq \emptyset$ ) do
     $u \leftarrow \text{ExtractMin}(Q)$           ** $s$  dequeued first
    foreach  $v$  neighbor of  $u$  do
      if ( $v.dist > u.dist + w(u, v)$ ) then
         $v.dist \leftarrow u.dist + w(u, v)$   **a Relax() call
         $v.predec \leftarrow u$ 
        decrease-key( $Q, v, v.dist$ )


```

Dijkstra's SSSP algorithm — an example:

▶ Input graph G :




▶ $\text{dijkstra}(G, 1)$:



▶ $\text{dijkstra}(G, 1)$ trace:

v	1	2	3	4	5	6	7
$v.\text{dist}/v.\text{predec}$	0/NIL	∞ /NIL	∞ /NIL	∞ /NIL	∞ /NIL	∞ /NIL	∞ /NIL
1 dequeued	0/NIL	9/1	∞ /NIL	∞ /NIL	3/1	12/1	∞ /NIL
5 dequeued	0/NIL	9/1	11/5	∞ /NIL	3/1	12/1	∞ /NIL
2 dequeued	0/NIL	9/1	10/2	∞ /NIL	3/1	12/1	19/2
3 dequeued	0/NIL	9/1	10/2	25/3	3/1	12/1	19/2
6 dequeued	0/NIL	9/1	10/2	25/3	3/1	12/1	19/2
7 dequeued	0/NIL	9/1	10/2	23/7	3/1	12/1	19/2



$$S = \{ \} \quad \{A, B, C, D, E\} = Q$$

$$\{A\} \quad \textcircled{0} \quad \infty \quad \infty \quad \infty$$

$$\{A, C\} \quad 0 \quad 10 \quad \textcircled{3} \quad \infty \quad \infty$$

$$\{A, C, E\} \quad \textcircled{0} \quad 7 \quad 8 \quad 11 \quad \textcircled{5}$$

Dijkstra's SSSP algorithm — Correctness:

- ▶ LI: “at the start of each while-loop iteration $u.dist = d(s, u)$ for all $u \notin Q$.”
- ▶ Initialization: The statement is vacuously true when Q holds all nodes. The statement is clearly true for the first vertex taken out of Q : the source s .
- ▶ Termination: At the end of the while-loop, Q is empty, so we have found all shortest-paths distances from s .
- ▶ Maintenance:
 - ▶ Suppose LI holds at the beginning of the iteration. Denote u as the node $\text{Extract-Min}(Q)$ takes out.
 - ▶ ASOC $d(s, u) < u.dist$.
 - ▶ Look at a shortest-path $s \rightarrow u$. $s \notin Q$ but $u \in Q$. Let (x, y) be the first edge such that $x \notin Q$ but $y \in Q$. (y might be u , x might be s .)
 - ▶ First, $d(s, y) \leq d(s, u)$ (subpath optimality + non negative weights)
 - ▶ Second, when we took x out of Q we made sure $y.dist \leq x.dist + w(x, y)$
 - ▶ Since $x \notin Q$ in the beginning of the iteration, then $x.dist = d(s, x)$.
 - ▶ Thus $y.dist \leq d(s, x) + w(x, y) = d(s, y)$ (Again, subpath optimality)
 - ▶ Altogether: $y.dist = d(s, y) \leq d(s, u) < u.dist$
 - ▶ If $y = u$ — immediate contradiction.
If $y \neq u$ — then $\text{Extract-Min}(Q)$ should return y not u . Contradiction in any case.

Dijkstra's SSSP algorithm — analysis:

- ▶ $|V| = n$ and $|E| = m$
- ▶ Running time:
 - ▶ `init()` — takes $O(n)$ time.
 - ▶ Initializing the priority-queue — $O(n)$
 - ▶ For each node, `ExtractMin` takes $O(\log(n))$ time.
 - ▶ ... and we search for all neighbors ($O(\deg(v))$ in the adjacency list model, $O(n)$ in the adjacency matrix model)
 - ▶ For every edge, we examine the edge at most twice (each time we take its endpoints from Q) and invoke `relax()` at most once, and decrease the key at most once.
So $O(\log(n))$ work per edge.
- ▶ In the adjacency-list model

$$O(n) + O(n) + O(n \log(n)) + O\left(\sum_v \deg(v)\right) + O(m \log(n)) = O((n+m) \log(n))$$

- ▶ In the adjacency-matrix model

$$O(n) + O(n) + O(n(\log(n) + n)) + O(m \log(n)) = O(n^2 + m \log(n))$$

- ▶ There exists a more refined implementation of the PQ (with Fibonacci heaps) that gives runtime $O(n \log(n) + m)$

Week 14: SSSP and APSP

Agenda:

- ▶ Bellman-Ford
- ▶ Floyd-Warshall

Reading:

- ▶ Textbook pages 643-663, 684-699

Bellman-Ford's SSSP algorithm for the general case:

- ▶ General case — edge weights **could be negative**
- ▶ Output:
 - Case 1. if there is a negative weight cycle, report it
 - Case 2. otherwise report all the $dist[u]$ values and the associated paths
- ▶ An key idea in the algorithm:
 - If there is no negative weight cycle reachable from s ,
 - every s -to- u shortest path contains at most $n - 1$ edges;
 - $\exists u$ such that s -to- u shortest path contains 1 edge;
 - ...
 - at termination: for every directed edge (u, v) there must be $d[v] \leq d[u] + w(u, v)$.
 - $d[v]$ can be reduced $n - 1$ times in order to reach value $dist[v]$, but no more.
- ▶ procedure `relax` (u, v)
 - if $d[v] > d[u] + w(u, v)$ then
 - $d[v] \leftarrow d[u] + w(u, v)$
 - $p[v] \leftarrow u$

- ▶ procedure bellman-ford(G, w, s) **** $G = (V, E)$**

```

for each  $v \in V(G)$  do           **initialization
     $d[v] \leftarrow \infty$ 
     $p[v] \leftarrow \text{NIL}$ 
 $d[s] \leftarrow 0$ 
for  $i \leftarrow 1$  to  $n - 1$        ** $n = |V(G)|$ 
    for each edge  $(u, v) \in E(G)$  do
        relax  $(u, v)$            **update  $d[v]$ 
for each edge  $(u, v) \in E(G)$  do
    if  $d[u] + w(u, v) < d[v]$  then **there is a negative cycle
        return FALSE
return TRUE

```
- ▶ Define $\delta(s, v)$: length of the shortest path from s to v
- ▶ Lemma 1: For all $v \in V$, at every iteration and over any sequence of relaxations we have $d[v] \geq \delta(s, v)$ and once $d[v]$ achieves value $\delta(s, v)$ it never changes.

- Proof: By induction on the number of relaxation steps. Base case $i = 0$ is trivial.

For I.S. consider the relaxation of an edge (u, v) and assume $d[v]$ changes:

$$\begin{aligned} d[v] &= d[u] + w(u, v) \\ &\geq \delta(s, u) + w(u, v) \\ &\geq \delta(s, v) \end{aligned}$$

Also note that $d[v]$ can only increase during a relax operation and since it is $\geq \delta(s, v)$ it doesn't change once it reaches this value.

- Lemma 2: Suppose that G has no negative cycle reachable from s . Then after $n - 1$ iterations $d[v] = \delta(s, v)$ for all $v \in V$.

Proof: Consider a vertex $v \in V$ and assume $s = v_0, v_1, v_2, \dots, v_k = v$ is a shortest path from s to v .

We prove by induction on i that after iteration $i \geq 0$: $d[v_i] = \delta(s, v_i)$.

Base: $i = 0$ is trivial

I.S.: Suppose that $d[v_{i-1}] = \delta(s, v_{i-1})$ and consider iteration i when we relax (v_{i-1}, v_i) . Then it can be seen that

$$d[v_i] = d[v_{i-1}] + w(v_{i-1}, v_i) = \delta(s, v_{i-1}) + w(v_{i-1}, v_i) = \delta(s, v_i).$$

- ▶ To complete the proof of correctness, we show that if there is a negative cycle then the algorithm detects it.
- ▶ suppose there is a negative cycle $c: v_0, v_1, v_2, \dots, v_k = v_0$:

$$\sum_{i=1}^k w(v_{i-1}, v_i) < 0$$

- ▶ By way of contradiction suppose the algorithm returns true, i.e. for all $1 \leq i \leq k$: $d[v_i] \leq d[v_{i-1}] + w(v_{i-1}, v_i)$. Thus:

$$\sum_{i=1}^k d[v_i] \leq \sum_{i=1}^k d[v_{i-1}] + \sum_{i=1}^k w(v_{i-1}, v_i)$$

- ▶ but $\sum_{i=1}^k d[v_i] = \sum_{i=1}^k d[v_{i-1}]$ since each vertex of the cycle appears exactly once in each sum.
Thus $0 \leq \sum_{i=1}^k w(v_{i-1}, v_i)$ which contradicts the assumption.

Floyd-Warshall's algorithm for All-Pairs-Shortest path:

- ▶ General case (weights can be negative but no negative cycle);
- ▶ Output: shortest path between “every” pair u, v of vertices.
- ▶ Idea: Use dynamic programming.
 Define $d[i, j, k]$ to be the length of shortest path from i to j for which all intermediate vertices are in $\{1, \dots, k\}$, for every $1 \leq i, j \leq n$ and $0 \leq k \leq n$.
- ▶ When $k = 0 \implies$ no intermediate vertex, so: $d[i, j, 0] = w(i, j)$.
- ▶ For general $k \geq 1$:
 - ▶ If the path does not contain k , inter. vertices only from $\{1, \dots, k-1\}$: $d[i, j, k-1]$.
 - ▶ If the path contains k , it has two parts: one goes from i to k with intermediate only from $\{1, \dots, k-1\}$, followed by a path from k to j using only from $\{1, \dots, k-1\}$: $d[i, k, k-1] + d[k, j, k-1]$.
- ▶ Recurrence:

$$d[i, j, k] = \min \begin{cases} w(i, j) & k = 0 \\ d[i, k, k-1] + d[k, j, k-1] & k \geq 1 \end{cases}$$

- ▶ We compute the table bottom-up, starting from smaller values of k to larger values.

Floyd-Warshall's algorithm:▶ ~~Pseudocode.~~

```

procedure Floyd-Warshall( $G$ )
  for  $i \leftarrow 1$  to  $n$  do
    for  $j \leftarrow 1$  to  $n$  do
       $d[i, j, 0] = w(i, j)$ 
       $b[i, j] = 0$       **  $b[i, j]$  keeps the break point vertex
  for  $k \leftarrow 1$  to  $n$  do
    for  $i \leftarrow 1$  to  $n$  do
      for  $j \leftarrow 1$  to  $n$  do
        if  $d[i, k, k-1] + d[k, j, k-1] < d[i, j, k-1]$  then
           $d[i, j, k] = d[i, k, k-1] + d[k, j, k-1]$ 
           $b[i, j] = k$ 
        else
           $d[i, j, k] = d[i, j, k-1]$ 

```

Floyd-Warshall's algorithm:

- ▶ To print the actual path between i and j :

```
procedure Print-FW( $i, j$ )  
  if  $b[i, j] = 0$  then  
    Print ‘‘edge’’,  $i$ , ‘‘to’’,  $j$   
  else  
    Print-FW( $i, b[i, j]$ )  
    Print-FW( $b[i, j], j$ )
```

- ▶ Running time: $\Theta(n^3)$, better than n run of Dijkstra (one for every vertex as a source).

